

# **Simple Neural Network in PyTorch**

# Why Google Colab??

---

- Google Colab provides a free cloud environment with a GPU (Graphic Processing Unit),
- It speeds up computations for deep learning tasks.
  - Runtime > Change runtime type > Hardware accelerator > GPU.
  - It will train the NN faster compared to using just the CPU.

## Note:

- ! tells the Colab environment to execute the command in the shell/terminal.
- Without the !, the command would be treated as a Python statement and would result in an error.

# Step 1: install PyTorch library (torch)

---

**!pip install torch torchvision**

**torchvision** adds additional features that are important if you are working with image-based machine learning projects. **Torchvision** provides

- Pre-trained Models: ResNet (Residual network), VGG (Visual Geometry Group (VGG))
  - These models can be used for image classification, object detection,  
Datasets: MNIST, CIFAR-10, ImageNet
  - These datasets come with utilities for downloading, loading, and preprocessing.
- Transforms: Resizing, cropping, normalization and converting to tensors.
  - Required to prepare image data before feeding it into a neural network.

## Step 2: Import the required libraries:

---

**import torch** #For creating and managing the neural network.

**import torch.nn as nn** #For defining the structure of the neural network.

**import torch.optim as optim** #For optimization (like gradient descent).

**import torchvision** # For working with standard datasets like MNIST

**import torchvision.transforms as transforms** # is used for image preprocessing

**import matplotlib.pyplot as plt** #For data visualization (create a variety of plots, such as line charts, bar charts, scatter plots, histograms)

## Step 3: Preparing the Dataset

**MNIST dataset** is used that has handwritten digits from 0 to 9 and each image in the dataset represents a single digit.

- Convert images to tensors (PyTorch's data format)
- Normalize the image pixel values to make training easier.
  - helps make training faster and more stable

## Step 3 Contd..

---

### Transformations:

- ToTensor(): Converts the image into a PyTorch tensor
  - Scales pixel values from [0, 255] to [0.0, 1.0].
- Normalize((0.5,), (0.5,)):
  - Normalizes pixel values to the range [-1, 1]  
$$\text{Normalized Pixel} = (\text{Pixel Value} - \text{Mean}) / \text{Std}$$
    - Center the data around 0.
    - To train the model faster.
    - Symmetric Activation Functions

### DataLoader:

- Organizes the dataset into batches for faster and more efficient training.
- train\_loader is for training, and test\_loader is for evaluation.

## Step 4: Defining the Neural Network

A neural network is a series of layers of mathematical operations. The simplest form is a feedforward neural network, where data flows in one direction

input → hidden layers → output

**self.fc1 = nn.Linear(28 \* 28, 128) # Input to hidden**

- It defines a fully connected layer **that maps 784 input features** (flattened MNIST pixels) to **128 output features** (hidden layer neurons).

**self.fc2 = nn.Linear(128, 64) # Hidden to hidden**

- It maps **128 features** and transforms it into another set of features **64 features**
  - ✓ It helps models to reduces computation

## Step 4 Contd..

**self.fc3 = nn.Linear(64, 10)      # Hidden to output**

- Output size = 10 (number of classes).
- In the MNIST dataset, there are 10 classes representing the digits 0 to 9.
- Each of the 10 output neurons corresponds to one class.

**self.relu = nn.ReLU()      # Define the Relu Activation function,**

- ReLU stands for Rectified Linear Unit, i.e.;  $\max(0, x)$
- You can use ReLU multiple times in your network's forward pass without redefining it



## Step 4 Contd..

`self.softmax = nn.LogSoftmax(dim=1)`

- It is used in the final output layer of classification models
- The Softmax function converts raw output values into probabilities for multi-class classification.
- `dim=1` means Softmax is applied across column-wise (classes) for each row (sample).
  - For each row, the sum of probabilities is 1.

`x = x.view(-1, 28 * 28)` #Flatten input, reshapes x into a 2D tensor with shape because layer in NN expect a 2D input

**Example:** A batch of size 16, containing MNIST images ([16, 1, 28, 28]). Here 16 is the total image, 1 is the number of channel, (28, 28) is height and weight  
For input ([16, 1, 28, 28]), output will be [16, 784]

## Step 5: Define the Loss Function and Optimizer

**model = SimpleNN()** # Contains layers, activation functions, and the forward pass logic.

**criterion = nn.NLLLoss()** # **Loss function, Negative Log-Likelihood Loss**

- Measures how well the model's predictions match the ground truth (actual labels).
- Guides the optimizer on how to adjust the model's parameters, e.g. weight
- Calculates the negative logarithm of the probability of the correct class.
- **Lower NLLLoss indicates a better model.**

## Step 5: Contd.

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

- This sets up the optimizer, which **updates the model's parameters** during training to **minimize the loss**.
- Adam (Adaptive Moment Estimation) is an optimization algorithm that adjusts the learning rate for each parameter. It is widely used due to its adaptive learning rate capabilities.
- A smaller learning rate (**lr=0.001**) ensures stability and allows the network to gradually improve its predictions.
- By combining **NLLLoss()** and **Adam**, the neural network can learn from its mistakes and improve its performance over time.

## **Step 6: Train the Model**

## **Step 7: Test the Model**

**Thank You**

The statement, "The MNIST dataset has 28x28 pixel grayscale images, flattened into 784 values," means the following:

1. **MNIST Dataset** MNIST is a dataset of handwritten digits (0–9). Each image in the dataset represents a single digit.

2. **28x28 Pixel Grayscale Images** Each image in the dataset is a grayscale image (not colored, so it has a single intensity channel).

The image dimensions are 28x28 pixels, meaning the image has: 28 rows and 28 columns of pixels.

Each pixel holds an intensity value that ranges from 0 (black) to 255 (white), with shades of gray in between.